

Machine Learning

Regression Methods 2

FAST

DISCOVERING
THE FUTURE

Topics of previous lectures

- ✓ Ingredients of Machine Learning
- ✓ Classification Basics, Basic Linear Classifier
- ✓ K-Nearest Neighbours Classifier
- ✓ Naive Bayes Classifier
- ✓ Linear and Quadratic Discriminant Analysis
- ✓ Support Vector Machines (SVM)
- ✓ Decision Trees
- ✓ Ensemble Methods (Bagging, Weighted Voting, Stacking)
- ✓ Linear Regression

Topics of today's lecture

- Regularization
- Maximum Likelihood Estimation
- Logistic Regression
- Softmax Classifier
- Support Vector Regression (SVR)
- Regression Trees
- Ensemble Methods

- The main idea of regularization in machine learning is to penalize complex models

- The main idea of regularization in machine learning is to penalize complex models
- If two models have a similar loss on the training data then the simpler tends to be better on test data

Regularization

- The main idea of regularization in machine learning is to penalize complex models
- If two models have a similar loss on the training data then the simpler tends to be better on test data
- Simpler models have less capacity to overfit the noise

Ridge regression

- Ridge regression is a regularisation (also known as L2-norm regularization) method for linear regression

Ridge regression

- Ridge regression is a regularisation (also known as L2-norm regularization) method for linear regression
- It regularises the task by penalising vector length:

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

Ridge regression

- Ridge regression is a regularisation (also known as L2-norm regularization) method for linear regression
- It regularises the task by penalising vector length:

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

- Here λ is a regularisation parameter, which we can choose by cross-validation:
 - Higher λ means stronger regularization
 - If $\lambda = 0$ then we are back at OLS

Calculating the partial derivatives

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

- Based on properties of matrix derivatives:

$$\frac{\partial \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \right)}{\partial \mathbf{w}} =$$

Calculating the partial derivatives

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

- Based on properties of matrix derivatives:

$$\frac{\partial \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \right)}{\partial \mathbf{w}} = -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + 2\lambda \mathbf{w} = 0$$

Calculating the partial derivatives

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

- Based on properties of matrix derivatives:

$$\frac{\partial \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \right)}{\partial \mathbf{w}} = -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + 2\lambda \mathbf{w} = 0$$
$$\mathbf{X}^T \mathbf{X} \mathbf{w} + \lambda \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

Calculating the partial derivatives

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

- Based on properties of matrix derivatives:

$$\frac{\partial \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \right)}{\partial \mathbf{w}} = -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + 2\lambda \mathbf{w} = 0$$
$$\mathbf{X}^T \mathbf{X} \mathbf{w} + \lambda \mathbf{w} = \mathbf{X}^T \mathbf{y}$$
$$(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

Calculating the partial derivatives

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

- Based on properties of matrix derivatives:

$$\frac{\partial \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \right)}{\partial \mathbf{w}} = -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + 2\lambda \mathbf{w} = 0$$

$$\mathbf{X}^T \mathbf{X} \mathbf{w} + \lambda \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{y})$$

Calculating the partial derivatives

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

- Based on properties of matrix derivatives:

$$\frac{\partial \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \right)}{\partial \mathbf{w}} = -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + 2\lambda \mathbf{w} = 0$$

$$\mathbf{X}^T \mathbf{X} \mathbf{w} + \lambda \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{y})$$

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{y})$$

Calculating the partial derivatives

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \|\mathbf{w}\|^2 \right)$$

- Based on properties of matrix derivatives:

$$\frac{\partial \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \right)}{\partial \mathbf{w}} = -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + 2\lambda \mathbf{w} = 0$$

$$\mathbf{X}^T \mathbf{X} \mathbf{w} + \lambda \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{y})$$

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{y})$$

- For ridge regression there exists a **closed-form solution** (i.e. can be explicitly calculated without numerical optimization):

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} (\mathbf{X}^T \mathbf{y})$$

Ridge regression

Pro: effective in reducing overfitting compared to OLS

Ridge regression

Pro: effective in reducing overfitting compared to OLS

Con: all coefficients are still non-zero, even if the true model is very sparse
(very few non-zero coefficients)

Lasso regression

- Lasso regression is another regularisation (also known as L1-norm regularization) method for linear regression

- Lasso regression is another regularisation (also known as L1-norm regularization) method for linear regression
- It regularises the task by penalising the sum of absolute values of weights:

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \sum_{i=1}^d |w_i| \right)$$

- Lasso regression is another regularisation (also known as L1-norm regularization) method for linear regression
- It regularises the task by penalising the sum of absolute values of weights:

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \left((\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \sum_{i=1}^d |w_i| \right)$$

- No closed-form solution exists, must be optimised numerically

Lasso vs Ridge regression

- Lasso can produce sparse solutions:
 - Many coefficients are exactly zero
 - How many? Depends on λ

Lasso vs Ridge regression

- Lasso can produce sparse solutions:
 - Many coefficients are exactly zero
 - How many? Depends on λ
- This is beneficial for two reasons:
 - The learned model becomes easier to interpret
 - If the true model is sparse, then lasso tends to give lower test MSE than ridge regression

Lasso vs Ridge regression

- Lasso can produce sparse solutions:
 - Many coefficients are exactly zero
 - How many? Depends on λ
- This is beneficial for two reasons:
 - The learned model becomes easier to interpret
 - If the true model is sparse, then lasso tends to give lower test MSE than ridge regression
- The parameter can be learned by using a hold-out validation dataset or K-fold cross-validation

Pro: if the true model is sparse then outperforms ridge regression and OLS

Lasso regression

- Pro:** if the true model is sparse then outperforms ridge regression and OLS
- Con:** if several features are highly correlated then tends to put high weight on an arbitrary one of those and zero to others

Logistic Regression

- Despite the name, **logistic regression is a classification method**, or it can be viewed as a regression method restricted into $[0, 1]$ interval

Logistic Regression

- Despite the name, **logistic regression is a classification method**, or it can be viewed as a regression method restricted into $[0, 1]$ interval
- We have used **Linear Regression** when y was continuous, but will it be easily used when y is binary?

Logistic Regression

- Despite the name, **logistic regression is a classification method**, or it can be viewed as a regression method restricted into $[0, 1]$ interval
- We have used **Linear Regression** when y was continuous, but will it be easily used when y is binary?
- We need a way to transform the predictions, so that we can get binary labels, instead of continuous quantities

Logistic Regression

- Despite the name, **logistic regression is a classification method**, or it can be viewed as a regression method restricted into $[0, 1]$ interval
- We have used **Linear Regression** when y was continuous, but will it be easily used when y is binary?
- We need a way to transform the predictions, so that we can get binary labels, instead of continuous quantities
- We can use the logistic sigmoid function on the predictions!

$$\sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$


Logistic Regression

- Despite the name, **logistic regression is a classification method**, or it can be viewed as a regression method restricted into $[0, 1]$ interval
- We have used **Linear Regression** when y was continuous, but will it be easily used when y is binary?
- We need a way to transform the predictions, so that we can get binary labels, instead of continuous quantities
- We can use the logistic sigmoid function on the predictions!

$$\sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

- In Logistic Regression the parameters are usually fit by **Maximum Likelihood Estimation (MLE)** method

Motivation for MLE: Population vs Sample



Maximum Likelihood Estimation (MLE)

Definition (Likelihood function)

Let $f(x_1, \dots, x_n; \theta)$, $\theta \in \mathbb{R}^k$ be the joint PMF or PDF of random variables $X_1, \dots, X_n$ with sample values $x_1, \dots, x_n$. The **likelihood function** of the sample is given by

$$L(\theta; x_1, \dots, x_n) = L(\theta) = f(x_1, \dots, x_n; \theta)$$

If $X_1, \dots, X_n$ are independent and identically distributed (i.i.d.)

Maximum Likelihood Estimation (MLE)

Definition (Likelihood function)

Let $f(x_1, \dots, x_n; \theta)$, $\theta \in \mathbb{R}^k$ be the joint PMF or PDF of random variables $X_1, \dots, X_n$ with sample values $x_1, \dots, x_n$. The **likelihood function** of the sample is given by

$$L(\theta; x_1, \dots, x_n) = L(\theta) = f(x_1, \dots, x_n; \theta)$$

If $X_1, \dots, X_n$ are independent and identically distributed (i.i.d.)

- discrete random variable with PMF $p(x, \theta)$, then

$$L(\theta) = P(X_1 = x_1, \dots, X_n = x_n) = \prod_{i=1}^n P(X_i = x_i) = \prod_{i=1}^n p(x_i, \theta)$$

Maximum Likelihood Estimation (MLE)

Definition (Likelihood function)

Let $f(x_1, \dots, x_n; \theta)$, $\theta \in \mathbb{R}^k$ be the joint PMF or PDF of random variables $X_1, \dots, X_n$ with sample values $x_1, \dots, x_n$. The **likelihood function** of the sample is given by

$$L(\theta; x_1, \dots, x_n) = L(\theta) = f(x_1, \dots, x_n; \theta)$$

If $X_1, \dots, X_n$ are independent and identically distributed (i.i.d.)

- discrete random variable with PMF $p(x, \theta)$, then

$$L(\theta) = P(X_1 = x_1, \dots, X_n = x_n) = \prod_{i=1}^n P(X_i = x_i) = \prod_{i=1}^n p(x_i, \theta)$$

- continuous r.v. with density $f(x, \theta)$, then similarly

$$L(\theta) = \prod_{i=1}^n f(x_i, \theta)$$

Maximum Likelihood Estimation (MLE)

Definition (Maximum Likelihood Estimators)

The maximum likelihood estimators (MLEs) are those values of the parameters that maximize the likelihood function with respect to the parameter θ . That is,

$$\hat{\theta} = \underset{\theta}{\operatorname{argmax}} L(\theta; x_1, \dots, x_n)$$

Maximum likelihood estimates give the parameter values for which the observed sample is **most likely** to have been generated.

Example

- Suppose the data $x_1, x_2, \dots, x_n$ is drawn independently from a normal distribution $N(\mu, \sigma^2)$ with unknown μ and σ

Example

- Suppose the data $x_1, x_2, \dots, x_n$ is drawn independently from a normal distribution $N(\mu, \sigma^2)$ with unknown μ and σ
- We want to estimate these unknown parameters from the data

Example

- Suppose the data $x_1, x_2, \dots, x_n$ is drawn independently from a normal distribution $N(\mu, \sigma^2)$ with unknown μ and σ
- We want to estimate these unknown parameters from the data
- For which values of μ and σ it is **most likely** that our data comes from the corresponding normal distribution?

$$\begin{aligned}\hat{\mu}, \hat{\sigma} &= \operatorname{argmax}_{\mu, \sigma} L(\mu, \sigma) = \prod_{i=1}^n f(x_i; \mu, \sigma) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}} = \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma}\right)^n \exp\left(-\sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}\right)\end{aligned}$$

Example

- Suppose the data $x_1, x_2, \dots, x_n$ is drawn independently from a normal distribution $N(\mu, \sigma^2)$ with unknown μ and σ
- We want to estimate these unknown parameters from the data
- For which values of μ and σ it is **most likely** that our data comes from the corresponding normal distribution?

$$\begin{aligned}\hat{\mu}, \hat{\sigma} &= \operatorname{argmax}_{\mu, \sigma} L(\mu, \sigma) = \prod_{i=1}^n f(x_i; \mu, \sigma) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}} = \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma}\right)^n \exp\left(-\sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}\right)\end{aligned}$$

We can obtain a more convenient, but equivalent optimization problem if we consider the **log likelihood** function

Example

- Suppose the data $x_1, x_2, \dots, x_n$ is drawn independently from a normal distribution $N(\mu, \sigma^2)$ with unknown μ and σ
- We want to estimate these unknown parameters from the data
- For which values of μ and σ it is **most likely** that our data comes from the corresponding normal distribution?

$$\begin{aligned}\hat{\mu}, \hat{\sigma} &= \operatorname{argmax}_{\mu, \sigma} L(\mu, \sigma) = \prod_{i=1}^n f(x_i; \mu, \sigma) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}} = \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \right)^n \exp \left(- \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2} \right)\end{aligned}$$

We can obtain a more convenient, but equivalent optimization problem if we consider the **log likelihood** function

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

Example

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

Example

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

$$\frac{\partial \log L(\mu, \sigma)}{\partial \mu} =$$

Example

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

$$\frac{\partial \log L(\mu, \sigma)}{\partial \mu} = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = 0 \Rightarrow$$

Example

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

$$\frac{\partial \log L(\mu, \sigma)}{\partial \mu} = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = 0 \Rightarrow \sum_{i=1}^n x_i = n\mu \Rightarrow$$

Example

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

$$\frac{\partial \log L(\mu, \sigma)}{\partial \mu} = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = 0 \Rightarrow \sum_{i=1}^n x_i = n\mu \Rightarrow \hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

Example

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

$$\frac{\partial \log L(\mu, \sigma)}{\partial \mu} = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = 0 \Rightarrow \sum_{i=1}^n x_i = n\mu \Rightarrow \hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\frac{\partial \log L(\mu, \sigma)}{\partial \sigma} =$$

Example

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

$$\frac{\partial \log L(\mu, \sigma)}{\partial \mu} = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = 0 \Rightarrow \sum_{i=1}^n x_i = n\mu \Rightarrow \hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\frac{\partial \log L(\mu, \sigma)}{\partial \sigma} = \frac{n}{\sigma} + \sum_{i=1}^n \frac{(x_i - \mu)^2}{\sigma^3} = 0 \Rightarrow$$

Example

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

$$\frac{\partial \log L(\mu, \sigma)}{\partial \mu} = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = 0 \Rightarrow \sum_{i=1}^n x_i = n\mu \Rightarrow \hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\frac{\partial \log L(\mu, \sigma)}{\partial \sigma} = \frac{n}{\sigma} + \sum_{i=1}^n \frac{(x_i - \mu)^2}{\sigma^3} = 0 \Rightarrow \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

Example

$$\hat{\mu}, \hat{\sigma} = \underset{\mu, \sigma}{\operatorname{argmax}} \log L(\mu, \sigma) = -n \log(\sqrt{2\pi}) - n \log(\sigma) - \sum_{i=1}^n \frac{(x_i - \mu)^2}{2\sigma^2}$$

How can we find the optimal values $\hat{\mu}, \hat{\sigma}$?

$$\frac{\partial \log L(\mu, \sigma)}{\partial \mu} = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = 0 \Rightarrow \sum_{i=1}^n x_i = n\mu \Rightarrow \hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\frac{\partial \log L(\mu, \sigma)}{\partial \sigma} = \frac{n}{\sigma} + \sum_{i=1}^n \frac{(x_i - \mu)^2}{\sigma^3} = 0 \Rightarrow \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

In statistical terms, we call $\hat{\mu}$ and $\hat{\sigma}$ **point estimators** or **statistics**.

Definition (Point Estimator)

Let $x_1, x_2, \dots, x_n$ be a set of i.i.d. data points. A **point estimator** or **statistic** is any function of the data

$$\hat{\theta} = \hat{\theta}_n = g(x_1, x_2, \dots, x_n)$$

Definition (Point Estimator)

Let $x_1, x_2, \dots, x_n$ be a set of i.i.d. data points. A **point estimator** or **statistic** is any function of the data

$$\hat{\theta} = \hat{\theta}_n = g(x_1, x_2, \dots, x_n)$$

Since the data is drawn from a random process, any function of the data is random. Therefore $\hat{\theta}$ is a random variable.

Properties of Estimators

Definition (Point Estimator)

Let $x_1, x_2, \dots, x_n$ be a set of i.i.d. data points. A **point estimator** or **statistic** is any function of the data

$$\hat{\theta} = \hat{\theta}_n = g(x_1, x_2, \dots, x_n)$$

Since the data is drawn from a random process, any function of the data is random. Therefore $\hat{\theta}$ is a random variable.

Definition (Bias)

The bias of an estimator is defined as

$$\text{bias}(\hat{\theta}_n) = \mathbb{E}(\hat{\theta}_n) - \theta,$$

where the expectation is taken over the data and θ is the true underlying value used to define the data-generating distribution.

Definition (Unbiased Estimator)

An estimator $\hat{\theta}_n$ is said to be **unbiased** if $\text{bias}(\hat{\theta}_n) = 0$, which implies that $\mathbb{E}(\hat{\theta}_n) = \theta$.

Properties of Estimators

Definition (Unbiased Estimator)

An estimator $\hat{\theta}_n$ is said to be **unbiased** if $bias(\hat{\theta}_n) = 0$, which implies that $\mathbb{E}(\hat{\theta}_n) = \theta$.

Definition (Asymptotically Unbiased)

An estimator $\hat{\theta}_n$ is said to be **asymptotically unbiased** if $\lim_{n \rightarrow \infty} bias(\hat{\theta}_n) = 0$, which implies that $\lim_{n \rightarrow \infty} \mathbb{E}(\hat{\theta}_n) = \theta$.

Properties of Estimators

Definition (Unbiased Estimator)

An estimator $\hat{\theta}_n$ is said to be **unbiased** if $bias(\hat{\theta}_n) = 0$, which implies that $\mathbb{E}(\hat{\theta}_n) = \theta$.

Definition (Asymptotically Unbiased)

An estimator $\hat{\theta}_n$ is said to be **asymptotically unbiased** if $\lim_{n \rightarrow \infty} bias(\hat{\theta}_n) = 0$, which implies that $\lim_{n \rightarrow \infty} \mathbb{E}(\hat{\theta}_n) = \theta$.

Definition (Consistency)

An estimator is weak consistent, if $\hat{\theta}_n \xrightarrow{p} \theta \quad n \rightarrow \infty$ and strong consistent, if $\hat{\theta}_n \xrightarrow{a.s.} \theta \quad n \rightarrow \infty$.

Example

The MLE estimator of the Gaussian mean parameter is unbiased. Suppose $x_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $i = 1, \dots, n$, then

$$\text{bias}(\hat{\mu}_n) = \mathbb{E}(\hat{\mu}_n) - \mu =$$

Example

The MLE estimator of the Gaussian mean parameter is unbiased. Suppose $x_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $i = 1, \dots, n$, then

$$bias(\hat{\mu}_n) = \mathbb{E}(\hat{\mu}_n) - \mu = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n x_i\right) - \mu =$$

Example

The MLE estimator of the Gaussian mean parameter is unbiased. Suppose $x_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $i = 1, \dots, n$, then

$$\text{bias}(\hat{\mu}_n) = \mathbb{E}(\hat{\mu}_n) - \mu = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n x_i\right) - \mu = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(x_i) - \mu = \mu - \mu = 0$$

Example

The MLE estimator of the Gaussian mean parameter is unbiased. Suppose $x_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $i = 1, \dots, n$, then

$$\text{bias}(\hat{\mu}_n) = \mathbb{E}(\hat{\mu}_n) - \mu = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n x_i\right) - \mu = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(x_i) - \mu = \mu - \mu = 0$$

The sample variance of a Gaussian distribution is biased, because

$$\text{bias}(\hat{\sigma}_n^2) = \mathbb{E}(\hat{\sigma}_n^2) - \sigma^2$$

Example

The MLE estimator of the Gaussian mean parameter is unbiased. Suppose $x_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $i = 1, \dots, n$, then

$$\text{bias}(\hat{\mu}_n) = \mathbb{E}(\hat{\mu}_n) - \mu = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n x_i\right) - \mu = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(x_i) - \mu = \mu - \mu = 0$$

The sample variance of a Gaussian distribution is biased, because

$$\begin{aligned} \text{bias}(\hat{\sigma}_n^2) &= \mathbb{E}(\hat{\sigma}_n^2) - \sigma^2 \\ \mathbb{E}(\hat{\sigma}_n^2) &= \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu}_n)^2\right) = \end{aligned}$$

Example

The MLE estimator of the Gaussian mean parameter is unbiased. Suppose $x_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $i = 1, \dots, n$, then

$$\text{bias}(\hat{\mu}_n) = \mathbb{E}(\hat{\mu}_n) - \mu = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n x_i\right) - \mu = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(x_i) - \mu = \mu - \mu = 0$$

The sample variance of a Gaussian distribution is biased, because

$$\begin{aligned} \text{bias}(\hat{\sigma}_n^2) &= \mathbb{E}(\hat{\sigma}_n^2) - \sigma^2 \\ \mathbb{E}(\hat{\sigma}_n^2) &= \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu}_n)^2\right) = \frac{n-1}{n} \sigma^2 (\text{show it!}) \Rightarrow \end{aligned}$$

Example

The MLE estimator of the Gaussian mean parameter is unbiased. Suppose $x_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $i = 1, \dots, n$, then

$$\text{bias}(\hat{\mu}_n) = \mathbb{E}(\hat{\mu}_n) - \mu = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n x_i\right) - \mu = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(x_i) - \mu = \mu - \mu = 0$$

The sample variance of a Gaussian distribution is biased, because

$$\begin{aligned}\text{bias}(\hat{\sigma}_n^2) &= \mathbb{E}(\hat{\sigma}_n^2) - \sigma^2 \\ \mathbb{E}(\hat{\sigma}_n^2) &= \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu}_n)^2\right) = \frac{n-1}{n} \sigma^2 (\text{show it!}) \Rightarrow \\ \Rightarrow \text{bias}(\hat{\sigma}_n^2) &= \frac{n-1}{n} \sigma^2 - \sigma^2 = -\frac{\sigma^2}{n}\end{aligned}$$

Example

The MLE estimator of the Gaussian mean parameter is unbiased. Suppose $x_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $i = 1, \dots, n$, then

$$\text{bias}(\hat{\mu}_n) = \mathbb{E}(\hat{\mu}_n) - \mu = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n x_i\right) - \mu = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(x_i) - \mu = \mu - \mu = 0$$

The sample variance of a Gaussian distribution is biased, because

$$\text{bias}(\hat{\sigma}_n^2) = \mathbb{E}(\hat{\sigma}_n^2) - \sigma^2$$

$$\mathbb{E}(\hat{\sigma}_n^2) = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu}_n)^2\right) = \frac{n-1}{n} \sigma^2 (\text{show it!}) \Rightarrow$$

$$\Rightarrow \text{bias}(\hat{\sigma}_n^2) = \frac{n-1}{n} \sigma^2 - \sigma^2 = -\frac{\sigma^2}{n}$$

The unbiased sample variance estimator is

$$\frac{1}{n-1} \sum_{i=1}^n (x_i - \hat{\mu}_n)^2.$$

MLE for Logistic Regression

- The likelihood function will be

$$L = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i},$$

where $p_i = P(y_i = \oplus | \mathbf{x}_i) = \sigma(\mathbf{w}^T \mathbf{x}_i)$

MLE for Logistic Regression

- The likelihood function will be

$$L = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i},$$

where $p_i = P(y_i = \oplus | \mathbf{x}_i) = \sigma(\mathbf{w}^T \mathbf{x}_i)$

- It is more convenient to work with the log likelihood function

$$\begin{aligned} \log(L) &= \sum_i \log(p_i^{y_i} (1 - p_i)^{1-y_i}) \\ &= \sum_i \left(\log(p_i^{y_i}) + \log((1 - p_i)^{1-y_i}) \right) = \\ &= \sum_i \left(y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \right) \end{aligned}$$

Does this look familiar?

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} = \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) =$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i \left(\frac{y_i}{\sigma(\mathbf{w}^T \mathbf{x}_i)} \frac{\partial}{\partial \mathbf{w}} (\sigma(\mathbf{w}^T \mathbf{x}_i)) \right) + \\ &+ \sum_i \left(\frac{1 - y_i}{1 - \sigma(\mathbf{w}^T \mathbf{x}_i)} \frac{\partial}{\partial \mathbf{w}} (1 - \sigma(\mathbf{w}^T \mathbf{x}_i)) \right) =\end{aligned}$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i \left(\frac{y_i}{\sigma(\mathbf{w}^T \mathbf{x}_i)} \frac{\partial}{\partial \mathbf{w}} (\sigma(\mathbf{w}^T \mathbf{x}_i)) \right) + \\ &+ \sum_i \left(\frac{1 - y_i}{1 - \sigma(\mathbf{w}^T \mathbf{x}_i)} \frac{\partial}{\partial \mathbf{w}} (1 - \sigma(\mathbf{w}^T \mathbf{x}_i)) \right) =\end{aligned}$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i \left(\frac{y_i}{\sigma(\mathbf{w}^T \mathbf{x}_i)} \frac{\partial}{\partial \mathbf{w}} (\sigma(\mathbf{w}^T \mathbf{x}_i)) \right) + \\ &+ \sum_i \left(\frac{1 - y_i}{1 - \sigma(\mathbf{w}^T \mathbf{x}_i)} \frac{\partial}{\partial \mathbf{w}} (1 - \sigma(\mathbf{w}^T \mathbf{x}_i)) \right) = \\ \text{the derivative of sigmoid is } \frac{d}{dz} \sigma(z) &= \sigma(z)(1 - \sigma(z))\end{aligned}$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i \left(\frac{y_i}{\sigma(\mathbf{w}^T \mathbf{x}_i)} \frac{\partial}{\partial \mathbf{w}} (\sigma(\mathbf{w}^T \mathbf{x}_i)) \right) + \\ &+ \sum_i \left(\frac{1 - y_i}{1 - \sigma(\mathbf{w}^T \mathbf{x}_i)} \frac{\partial}{\partial \mathbf{w}} (1 - \sigma(\mathbf{w}^T \mathbf{x}_i)) \right) =\end{aligned}$$

the derivative of sigmoid is $\frac{d}{dz} \sigma(z) = \sigma(z)(1 - \sigma(z))$

$$\begin{aligned}&= \sum_i \left(\frac{y_i}{\sigma(\mathbf{w}^T \mathbf{x}_i)} \sigma(\mathbf{w}^T \mathbf{x}_i)(1 - \sigma(\mathbf{w}^T \mathbf{x}_i)) \mathbf{x}_i \right) + \\ &+ \sum_i \left(-\frac{1 - y_i}{1 - \sigma(\mathbf{w}^T \mathbf{x}_i)} \sigma(\mathbf{w}^T \mathbf{x}_i)(1 - \sigma(\mathbf{w}^T \mathbf{x}_i)) \mathbf{x}_i \right)\end{aligned}$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i (y_i(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))\mathbf{x}_i) + \sum_i (-(1 - y_i)\sigma(\mathbf{w}^T \mathbf{x}_i)\mathbf{x}_i) =\end{aligned}$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i (y_i (1 - \sigma(\mathbf{w}^T \mathbf{x}_i)) \mathbf{x}_i) + \sum_i (-(1 - y_i) \sigma(\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i) = \\ &= \sum_i (\mathbf{x}_i (y_i - y_i \sigma(\mathbf{w}^T \mathbf{x}_i) - \sigma(\mathbf{w}^T \mathbf{x}_i) + y_i \sigma(\mathbf{w}^T \mathbf{x}_i))) =\end{aligned}$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i (y_i(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))\mathbf{x}_i) + \sum_i (-(1 - y_i)\sigma(\mathbf{w}^T \mathbf{x}_i)\mathbf{x}_i) = \\ &= \sum_i (\mathbf{x}_i(y_i - y_i\sigma(\mathbf{w}^T \mathbf{x}_i) - \sigma(\mathbf{w}^T \mathbf{x}_i) + y_i\sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i \mathbf{x}_i(y_i - \sigma(\mathbf{w}^T \mathbf{x}_i)) = 0\end{aligned}$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i (y_i(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))\mathbf{x}_i) + \sum_i (-(1 - y_i)\sigma(\mathbf{w}^T \mathbf{x}_i)\mathbf{x}_i) = \\ &= \sum_i (\mathbf{x}_i(y_i - y_i\sigma(\mathbf{w}^T \mathbf{x}_i) - \sigma(\mathbf{w}^T \mathbf{x}_i) + y_i\sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i \mathbf{x}_i(y_i - \sigma(\mathbf{w}^T \mathbf{x}_i)) = 0\end{aligned}$$

$$\sum_i \mathbf{x}_i y_i = \sum_i \mathbf{x}_i \sigma(\mathbf{w}^T \mathbf{x}_i)$$

MLE for Logistic Regression

Now that we have the log likelihood function, we can compute the partial derivative w.r.t. $\mathbf{w}$

$$\begin{aligned}\frac{\partial \log(L(\mathbf{w}))}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \sum_i (y_i \log(\sigma(\mathbf{w}^T \mathbf{x}_i)) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i (y_i(1 - \sigma(\mathbf{w}^T \mathbf{x}_i))\mathbf{x}_i) + \sum_i (-(1 - y_i)\sigma(\mathbf{w}^T \mathbf{x}_i)\mathbf{x}_i) = \\ &= \sum_i (\mathbf{x}_i(y_i - y_i\sigma(\mathbf{w}^T \mathbf{x}_i) - \sigma(\mathbf{w}^T \mathbf{x}_i) + y_i\sigma(\mathbf{w}^T \mathbf{x}_i))) = \\ &= \sum_i \mathbf{x}_i(y_i - \sigma(\mathbf{w}^T \mathbf{x}_i)) = 0\end{aligned}$$

$$\sum_i \mathbf{x}_i y_i = \sum_i \mathbf{x}_i \sigma(\mathbf{w}^T \mathbf{x}_i)$$

This system of equalities is hard to solve directly!

Logistic Regression Solution

- The above system of equalities is hard to solve directly

Logistic Regression Solution

- The above system of equalities is hard to solve directly
- Instead, it is easier to optimize

$$\operatorname{argmin}_{\mathbf{w}} - \sum_i \left(y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \right)$$

Logistic Regression Solution

- The above system of equalities is hard to solve directly
- Instead, it is easier to optimize

$$\operatorname{argmin}_{\mathbf{w}} - \sum_i \left(y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \right)$$

- This can be done with a numerical method such as gradient descent

Logistic Regression Solution

- The above system of equalities is hard to solve directly
- Instead, it is easier to optimize

$$\operatorname{argmin}_{\mathbf{w}} - \sum_i \left(y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \right)$$

- This can be done with a numerical method such as gradient descent
- The loss function that logistic regression minimizes is known as **cross entropy** (also known as log-loss or binomial deviance)

Softmax Classifier

- Softmax classifier is the multi-class generalization of Logistic Regression

Softmax Classifier

- Softmax classifier is the multi-class generalization of Logistic Regression
- In case of a C class classification task, we will have a weight matrix $\mathbf{W}$ of size $C \times (d + 1)$, instead of a $d + 1$ dimensional vector

Softmax Classifier

- Softmax classifier is the multi-class generalization of Logistic Regression
- In case of a C class classification task, we will have a weight matrix $\mathbf{W}$ of size $C \times (d + 1)$, instead of a $d + 1$ dimensional vector
- The softmax function is used on each of the prediction scores

$$f_j(\mathbf{z}) = \frac{e^{z_j}}{\sum_k e^{z_k}}, \quad \text{where} \quad z_j = \mathbf{W}_j \mathbf{x}^T, \mathbf{z} = \mathbf{W} \mathbf{x}^T$$

Softmax Classifier

- Softmax classifier is the multi-class generalization of Logistic Regression
- In case of a C class classification task, we will have a weight matrix $\mathbf{W}$ of size $C \times (d + 1)$, instead of a $d + 1$ dimensional vector
- The softmax function is used on each of the prediction scores

$$f_j(\mathbf{z}) = \frac{e^{z_j}}{\sum_k e^{z_k}}, \quad \text{where} \quad z_j = \mathbf{W}_j \mathbf{x}^T, \mathbf{z} = \mathbf{W} \mathbf{x}^T$$

- The loss function will be cross-entropy again

$$L(\mathbf{y}, f(\mathbf{x})) = - \sum_{j=1}^C y_j \log f_j(\mathbf{z}),$$

where $\mathbf{y}$ is a one-hot encoded vector for the ground truth class.

Softmax Classifier

- Softmax classifier is the multi-class generalization of Logistic Regression
- In case of a C class classification task, we will have a weight matrix $\mathbf{W}$ of size $C \times (d + 1)$, instead of a $d + 1$ dimensional vector
- The softmax function is used on each of the prediction scores

$$f_j(\mathbf{z}) = \frac{e^{z_j}}{\sum_k e^{z_k}}, \quad \text{where} \quad z_j = \mathbf{W}_j \mathbf{x}^T, \mathbf{z} = \mathbf{W} \mathbf{x}^T$$


- The loss function will be cross-entropy again

$$L(\mathbf{y}, f(\mathbf{x})) = - \sum_{j=1}^C y_j \log f_j(\mathbf{z}),$$

where $\mathbf{y}$ is a one-hot encoded vector for the ground truth class.

- The problem can be solved with gradient descent.

Support Vector Machine (SVM) for Classification



Hard-margin SVM for classification

$$\mathbf{w}^*, t^* = \operatorname{argmin}_{\mathbf{w}, t} \frac{1}{2} \|\mathbf{w}\|^2$$

subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1, \quad y_i \in \{-1, 1\}, \quad i = 1, \dots, n$

Hard-margin SVM for classification

$$\mathbf{w}^*, t^* = \operatorname{argmin}_{\mathbf{w}, t} \frac{1}{2} \|\mathbf{w}\|^2$$

subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1, \quad y_i \in \{-1, 1\}, \quad i = 1, \dots, n$

We can solve the equivalent dual problem

$$\alpha_1^*, \dots, \alpha_n^* = \operatorname{argmax}_{\alpha_1, \dots, \alpha_n} -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i$$

subject to $\alpha_i \geq 0, \quad i = 1, \dots, n$ and $\sum_{i=1}^n \alpha_i y_i = 0$

Hard-margin SVM for classification

$$\mathbf{w}^*, t^* = \operatorname{argmin}_{\mathbf{w}, t} \frac{1}{2} \|\mathbf{w}\|^2$$

subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1, \quad y_i \in \{-1, 1\}, \quad i = 1, \dots, n$

We can solve the equivalent dual problem

$$\alpha_1^*, \dots, \alpha_n^* = \operatorname{argmax}_{\alpha_1, \dots, \alpha_n} -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i$$

subject to $\alpha_i \geq 0, \quad i = 1, \dots, n$ and $\sum_{i=1}^n \alpha_i y_i = 0$

From the result we can calculate:

$$\mathbf{w}^* = \sum_{i=1}^n \alpha_i^* y_i \mathbf{x}_i \quad t^* = \mathbf{w}^* \cdot \mathbf{x}_i - y_i,$$

where $\mathbf{x}_i$ is a support vector and α_i is its weight.

Soft-margin SVM for classification

$$\mathbf{w}^*, t^*, \xi_i^* = \operatorname{argmin}_{\mathbf{w}, t, \xi_i} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$

subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad i = 1, \dots, n$

Soft-margin SVM for classification

$$\mathbf{w}^*, t^*, \xi_i^* = \underset{\mathbf{w}, t, \xi_i}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$

$$\text{subject to } y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad i = 1, \dots, n$$

We can solve the equivalent dual problem

$$\alpha_1^*, \dots, \alpha_n^* = \underset{\alpha_1, \dots, \alpha_n}{\operatorname{argmax}} -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i$$

$$\text{subject to } 0 \leq \alpha_i \leq C, \quad i = 1, \dots, n \text{ and } \sum_{i=1}^n \alpha_i y_i = 0$$

Soft-margin SVM for classification

$$\mathbf{w}^*, t^*, \xi_i^* = \underset{\mathbf{w}, t, \xi_i}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$

$$\text{subject to } y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad i = 1, \dots, n$$

We can solve the equivalent dual problem

$$\alpha_1^*, \dots, \alpha_n^* = \underset{\alpha_1, \dots, \alpha_n}{\operatorname{argmax}} -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i$$


$$\text{subject to } 0 \leq \alpha_i \leq C, \quad i = 1, \dots, n \text{ and } \sum_{i=1}^n \alpha_i y_i = 0$$

From the result we can calculate:

$$\mathbf{w}^* = \sum_{i=1}^n \alpha_i^* y_i \mathbf{x}_i \quad t^* = \mathbf{w}^* \cdot \mathbf{x}_i - y_i,$$

where $\mathbf{x}_i$ is a support vector and α_i is its weight.

SVM for Regression



Support Vector Regression

$$\min_{\mathbf{w}, t, \xi_i, \xi'_i} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n (\xi_i + \xi'_i)$$

$$\text{subject to } y_i - \mathbf{w} \cdot \mathbf{x}_i - t \leq \epsilon + \xi_i, \quad \xi_i \geq 0$$

$$\mathbf{w} \cdot \mathbf{x}_i + t - y_i \leq \epsilon + \xi'_i, \quad \xi'_i \geq 0, \quad i = 1, \dots, n$$

Support Vector Regression

$$\min_{\mathbf{w}, t, \xi_i, \xi'_i} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n (\xi_i + \xi'_i)$$

$$\text{subject to } y_i - \mathbf{w} \cdot \mathbf{x}_i - t \leq \epsilon + \xi_i, \quad \xi_i \geq 0$$

$$\mathbf{w} \cdot \mathbf{x}_i + t - y_i \leq \epsilon + \xi'_i, \quad \xi'_i \geq 0, \quad i = 1, \dots, n$$

We can solve the equivalent dual problem

$$\max_{\boldsymbol{\alpha}; \boldsymbol{\alpha}'} -\frac{1}{2}(\boldsymbol{\alpha} - \boldsymbol{\alpha}')^T Q(\boldsymbol{\alpha} - \boldsymbol{\alpha}') - \epsilon e^T(\boldsymbol{\alpha} + \boldsymbol{\alpha}') + y^T(\boldsymbol{\alpha} - \boldsymbol{\alpha}')$$

$$\text{subject to } 0 \leq \alpha_i, \alpha'_i \leq C, \quad i = 1, \dots, n \text{ and } e^T(\boldsymbol{\alpha} - \boldsymbol{\alpha}') = 0$$

where $Q = \mathbf{x}\mathbf{x}^T$ and $e = (1, \dots, 1)^T$

Support Vector Regression

$$\min_{\mathbf{w}, t, \xi_i, \xi'_i} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n (\xi_i + \xi'_i)$$

$$\text{subject to } y_i - \mathbf{w} \cdot \mathbf{x}_i - t \leq \epsilon + \xi_i, \quad \xi_i \geq 0$$

$$\mathbf{w} \cdot \mathbf{x}_i + t - y_i \leq \epsilon + \xi'_i, \quad \xi'_i \geq 0, \quad i = 1, \dots, n$$

We can solve the equivalent dual problem

$$\max_{\alpha; \alpha'} -\frac{1}{2} (\alpha - \alpha')^T Q (\alpha - \alpha') - \epsilon e^T (\alpha + \alpha') + y^T (\alpha - \alpha')$$

$$\text{subject to } 0 \leq \alpha_i, \alpha'_i \leq C, \quad i = 1, \dots, n \text{ and } e^T (\alpha - \alpha') = 0$$

where $Q = \mathbf{x}\mathbf{x}^T$ and $e = (1, \dots, 1)^T$

From this we can calculate the prediction: $\hat{f}(\mathbf{x}) = \mathbf{w}^* \cdot \mathbf{x} + t$, with

$$\mathbf{w}^* = \sum_{i=1}^n (\alpha_i - \alpha'_i) \mathbf{x}_i, \quad t^* = y_i - \epsilon - \mathbf{w}^* \cdot \mathbf{x}_i$$

Support Vector Regression

Can we apply kernels in this case?

$$\min_{\mathbf{w}, t, \xi_i, \xi'_i} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n (\xi_i + \xi'_i)$$

$$\text{subject to } y_i - \mathbf{w} \cdot \mathbf{x}_i - t \leq \epsilon + \xi_i, \quad \xi_i \geq 0$$

$$\mathbf{w} \cdot \mathbf{x}_i + t - y_i \leq \epsilon + \xi'_i, \quad \xi'_i \geq 0, \quad i = 1, \dots, n$$

We can solve the equivalent dual problem

$$\max_{\alpha; \alpha'} -\frac{1}{2} (\alpha - \alpha')^T Q (\alpha - \alpha') - \epsilon e^T (\alpha + \alpha') + y^T (\alpha - \alpha')$$

$$\text{subject to } 0 \leq \alpha_i, \alpha'_i \leq C, \quad i = 1, \dots, n \text{ and } e^T (\alpha - \alpha') = 0$$

where $Q = \mathbf{xx}^T$ and $e = (1, \dots, 1)^T$

From this we can calculate the prediction: $\hat{f}(\mathbf{x}) = \mathbf{w}^* \cdot \mathbf{x} + t$, with

$$\mathbf{w}^* = \sum_{i=1}^n (\alpha_i - \alpha'_i) \mathbf{x}_i, \quad t^* = y_i - \epsilon - \mathbf{w}^* \cdot \mathbf{x}_i$$

- Fitting (soft-margin kernel SVR):

$$\max_{\alpha; \alpha'} -\frac{1}{2}(\alpha - \alpha')^T Q(\alpha - \alpha') + \epsilon e^T(\alpha + \alpha') - y^T(\alpha - \alpha')$$

subject to $0 \leq \alpha_i, \alpha'_i \leq C, \quad i = 1, \dots, n$ and $e^T(\alpha - \alpha') = 0$

where $Q = \kappa(\mathbf{x}_i, \mathbf{x}_j)$ and $e = (1, \dots, 1)^T$

$$t^* = y_j - \epsilon - \sum_{i=1}^n (\alpha_i - \alpha'_i) \kappa(\mathbf{x}_i, \mathbf{x}_j),$$

where $\mathbf{x}_j$ is a support vector of target value y_j .

- Fitting (soft-margin kernel SVR):

$$\max_{\alpha; \alpha'} -\frac{1}{2}(\alpha - \alpha')^T Q(\alpha - \alpha') + \epsilon e^T(\alpha + \alpha') - y^T(\alpha - \alpha')$$

subject to $0 \leq \alpha_i, \alpha'_i \leq C, \quad i = 1, \dots, n$ and $e^T(\alpha - \alpha') = 0$

where $Q = \kappa(\mathbf{x}_i, \mathbf{x}_j)$ and $e = (1, \dots, 1)^T$

$$t^* = y_j - \epsilon - \sum_{i=1}^n (\alpha_i - \alpha'_i) \kappa(\mathbf{x}_i, \mathbf{x}_j),$$

where $\mathbf{x}_j$ is a support vector of target value y_j .

- Prediction:

$$\hat{f}(\mathbf{x}) = \sum_{i=1}^n (\alpha_i - \alpha'_i) \kappa(\mathbf{x}_i, \mathbf{x}) + t^*$$

Regression Trees

- In order to adapt decision trees to regression problem, we need to think about a suitable impurity function

Regression Trees

- In order to adapt decision trees to regression problem, we need to think about a suitable impurity function
- We can use the **weighted average variance** of the target values as the measure of impurity

Regression Trees

- In order to adapt decision trees to regression problem, we need to think about a suitable impurity function
- We can use the **weighted average variance** of the target values as the measure of impurity
- If a split partitions the set of target values Y into mutually exclusive sets $\{Y_1, \dots, Y_l\}$, the weighted average variance is then

$$\mathbb{V}ar(\{Y_1, \dots, Y_l\}) = \sum_{j=1}^l \frac{|Y_j|}{|Y|} \mathbb{V}ar(Y_j) =$$

Regression Trees

- In order to adapt decision trees to regression problem, we need to think about a suitable impurity function
- We can use the **weighted average variance** of the target values as the measure of impurity
- If a split partitions the set of target values Y into mutually exclusive sets $\{Y_1, \dots, Y_l\}$, the weighted average variance is then

$$\mathbb{V}ar(\{Y_1, \dots, Y_l\}) = \sum_{j=1}^l \frac{|Y_j|}{|Y|} \mathbb{V}ar(Y_j) = \sum_{j=1}^l \frac{|Y_j|}{|Y|} \frac{1}{|Y_j|} \sum_{y \in |Y_j|} (y - \bar{y})^2 =$$

Regression Trees

- In order to adapt decision trees to regression problem, we need to think about a suitable impurity function
- We can use the **weighted average variance** of the target values as the measure of impurity
- If a split partitions the set of target values Y into mutually exclusive sets $\{Y_1, \dots, Y_l\}$, the weighted average variance is then

$$\begin{aligned}\mathbb{V}ar(\{Y_1, \dots, Y_l\}) &= \sum_{j=1}^l \frac{|Y_j|}{|Y|} \mathbb{V}ar(Y_j) = \sum_{j=1}^l \frac{|Y_j|}{|Y|} \frac{1}{|Y_j|} \sum_{y \in |Y_j|} (y - \bar{y})^2 = \\ &= \sum_{j=1}^l \left(\frac{|Y_j|}{|Y|} \frac{1}{|Y_j|} \sum_{y \in |Y_j|} y^2 - \bar{y}^2 \right) =\end{aligned}$$

Regression Trees

- In order to adapt decision trees to regression problem, we need to think about a suitable impurity function
- We can use the **weighted average variance** of the target values as the measure of impurity
- If a split partitions the set of target values Y into mutually exclusive sets $\{Y_1, \dots, Y_l\}$, the weighted average variance is then

$$\begin{aligned}\mathbb{V}ar(\{Y_1, \dots, Y_l\}) &= \sum_{j=1}^l \frac{|Y_j|}{|Y|} \mathbb{V}ar(Y_j) = \sum_{j=1}^l \frac{|Y_j|}{|Y|} \frac{1}{|Y_j|} \sum_{y \in |Y_j|} (y - \bar{y})^2 = \\ &= \sum_{j=1}^l \left(\frac{|Y_j|}{|Y|} \frac{1}{|Y_j|} \sum_{y \in |Y_j|} y^2 - \bar{y}^2 \right) = \\ &= \frac{1}{|Y|} \sum_{y \in |Y_j|} y^2 - \sum_{j=1}^l \frac{|Y_j|}{|Y|} \bar{y}^2\end{aligned}$$

- A split is considered good if it maximizes the impurity gain:

$$\text{Imp}(D) - \text{Imp}(\{D_1, \dots, D_l\})$$

the overall impurity of child nodes should be maximally smaller than the impurity of the parent D

- A split is considered good if it maximizes the impurity gain:

$$Imp(D) - Imp(\{D_1, \dots, D_l\})$$

the overall impurity of child nodes should be maximally smaller than the impurity of the parent D

- In case of regression we take $Imp = Var$ and we base our decision on **variance reduction**

- A split is considered good if it maximizes the impurity gain:

$$Imp(D) - Imp(\{D_1, \dots, D_l\})$$


the overall impurity of child nodes should be maximally smaller than the impurity of the parent D

- In case of regression we take $Imp = Var$ and we base our decision on **variance reduction**
- Each leaf returns the average target value of the segment

- In case of **bagging**, we take the **average prediction** across the individual models

Ensemble Methods

- In case of **bagging**, we take the **average prediction** across the individual models



- In case of **Random Forest**, we grow a forest of **regression trees** and take the average prediction across the trees. The number of features in each tree is usually taken as $F = \frac{P}{3}$, where P is the number of total features in the dataset

What have we learned today?

- ✓ Regularization
- ✓ Maximum Likelihood Estimation
- ✓ Logistic Regression
- ✓ Softmax Classifier
- ✓ Support Vector Regression (SVR)
- ✓ Regression Trees
- ✓ Ensemble Methods